

# CSE519\_Homework2\_Final

September 25, 2018

## 0.1 Imports

```
In [0]: import numpy as np           #For NumPy arrays and mathematical functions
import pandas as pd                 #For operations using DataFrames
import matplotlib.pyplot as plt     #For plotting
import seaborn as sns               #For plotting
import time
import os
```

## 0.2 Mounting Google Drive on Colab

```
In [2]: from google.colab import drive
drive.mount('gdrive')
```

Mounted at gdrive

Creating File and Directory paths

```
In [0]: dir_path = "gdrive/My Drive/Colab Notebooks/Datasets/NYC/"
file_name = "train.feather"
```

### 0.2.1 Installing Feather locally for reading

```
In [0]: !pip install -U feather-format
```

## 0.3 Reading the dataset

```
In [4]: print("Reading Feather file...")
df = pd.read_feather(dir_path+file_name)
print("Reading completed")
%time
```

Reading Feather file...

```
/usr/local/lib/python3.6/dist-packages/pandas/io/feather_format.py:112: FutureWarning: `nthreads`
return feather.read_dataframe(path, nthreads=nthreads)
```

```
Reading completed
CPU times: user 3 ȳs, sys: 2 ȳs, total: 5 ȳs
Wall time: 14.5 ȳs
```

### 0.3.1 Now let's examine the properties of dataset

```
In [0]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 55423856 entries, 0 to 55423855
Data columns (total 7 columns):
fare_amount           float32
pickup_datetime       datetime64[ns, UTC]
pickup_longitude       float32
pickup_latitude       float32
dropoff_longitude     float32
dropoff_latitude       float32
passenger_count       uint8
dtypes: datetime64[ns, UTC](1), float32(5), uint8(1)
memory usage: 1.5 GB
```

```
In [6]: df.head()
```

```
Out[6]:
```

	fare_amount	pickup_datetime	pickup_longitude	pickup_latitude \		dropoff_longitude	dropoff_latitude	passenger_count
0	4.5	2009-06-15 17:26:00+00:00	-73.844315	40.721317		-73.841614	40.712276	1
1	16.9	2010-01-05 16:52:00+00:00	-74.016045	40.711304		-73.979271	40.782005	1
2	5.7	2011-08-18 00:35:00+00:00	-73.982735	40.761269		-73.991241	40.750561	2
3	7.7	2012-04-21 04:30:00+00:00	-73.987129	40.733143		-73.991570	40.758091	1
4	5.3	2010-03-09 07:51:00+00:00	-73.968094	40.768009		-73.956657	40.783764	1

```
In [7]: df.describe()
```

```
Out[7]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude \	
count	5.542386e+07	5.542386e+07	5.542386e+07	5.542348e+07	
mean	1.134503e+01	-7.250972e+01	3.991985e+01	-7.251106e+01	
std	2.071083e+01	1.284888e+01	9.642353e+00	1.278220e+01	
min	-3.000000e+02	-3.442060e+03	-3.492264e+03	-3.442025e+03	
25%	6.000000e+00	-7.399207e+01	4.073493e+01	-7.399140e+01	
50%	8.500000e+00	-7.398180e+01	4.075265e+01	-7.398015e+01	
75%	1.250000e+01	-7.396708e+01	4.076713e+01	-7.396368e+01	

max	9.396336e+04	3.457626e+03	3.408790e+03	3.457622e+03
-----	--------------	--------------	--------------	--------------

	dropoff_latitude	passenger_count
count	5.542348e+07	5.542386e+07
mean	3.992067e+01	1.685380e+00
std	9.633346e+00	1.327664e+00
min	-3.547887e+03	0.000000e+00
25%	4.073403e+01	1.000000e+00
50%	4.075316e+01	1.000000e+00
75%	4.076810e+01	2.000000e+00
max	3.537133e+03	2.080000e+02

As we can see from the above description, our dataset contains a lot of anomalies and outliers. If a model is built without cleaning, it will be far from accurate. To build a good model, we will have to clean the data and add additional features

## 0.4 Data Cleaning and Exploration

### 1. Perform Initial Operations Check shape of data

```
In [0]: df.shape
```

```
Out[0]: (55423856, 7)
```

Check rows having NaN values

```
In [5]: df.isnull().sum().sort_values(ascending=False)
```

```
Out[5]: dropoff_latitude    376
dropoff_longitude    376
passenger_count         0
pickup_latitude        0
pickup_longitude        0
pickup_datetime        0
fare_amount            0
dtype: int64
```

Dropping rows having NaN values As only the first two columns contain NaN values, we choose a subset of column to check for dropping the rows and proceed with this to improve performance

```
In [0]: df = df.dropna(axis='index', how='any', subset=['dropoff_latitude', 'dropoff_longitude'])
```

Now, after the rows are dropped, the shape will change

```
In [9]: df.shape
```

```
Out[9]: (55423480, 7)
```

After having dropped the null rows, let us focus on making data good The data contains a lot of issues - 1. Fare Amount is negative 2. Latitude and Longitude 3. Passenger Count

**Step 1: Deleting rows having negative fare amount and fare above 95 dollars.** The first task towards cleaning data would be removing rows having negative fare amount.

```
In [0]: #Removing rows with negative fare
df = df[df.fare_amount > 0]
```


```
In [0]: df.shape
```

```
Out[0]: (55419646, 7)
```

Let's plot a histogram checking for the distribution of fare\_amount

```
In [0]: df[df.fare_amount<100].fare_amount.hist(bins=80)
plt.xlabel('Fare $ USD')
plt.ylabel('Frequency')
plt.title('Histogram')
```

```
Out[0]: Text(0.5,1,'Histogram')
```



This is a unimodal distribution. But even then, there exist fare amount greater than 150 dollars or so. A simple observation at the distribution reveals that the maximum fare amount lies between 5 to 15 dollars. In order to trim the outliers, let us check the quantiles.

```
In [0]: low = 0
high = .99953      #3.5-standard deviation, that means covering 99.953% observations
quantiles_df = df.quantile([low, high])
print(quantiles_df)
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
0.00000	0.010000	-3442.059570	-3492.263672	-3442.024658	
0.99953	95.330002	40.701997	41.012809	40.687964	

	dropoff_latitude	passenger_count
0.00000	-3547.886719	0.0
0.99953	40.989659	6.0

As we can see, for 3.5 sigma, fare\_amount -> 95 dollars. This means that the 99.953% of the observations lie below 95 dollars. **So it is safe to remove the outliers, that is, 0.047% observations.**

```
In [8]: #Removing rows having having fare amount less than 3.5 sigma
df = df[df.fare_amount <= 95]
df.shape
```

```
Out[8]: (55393378, 7)
```

```
In [0]: df.fare_amount.describe()
```

```
Out[0]: count      5.539338e+07
mean        1.128455e+01
std         9.371513e+00
min         1.000000e-02
25%         6.000000e+00
50%         8.500000e+00
75%         1.250000e+01
max         9.500000e+01
Name: fare_amount, dtype: float64
```

Even now, we can see that the minimum fare is \$0.01.

But looking online for the taxi fares during the years 2009-2015, we understand that the minimum fare was \$ 2. So we should only keep those rows for which the fare was above 2 dollars

```
In [9]: #Removing rows with fare amount less than USD 2.
df = df[df.fare_amount >= 2]
df.shape
```

```
Out[9]: (55392531, 7)
```

```
In [0]: low = 0
high = .9973          #3-standard deviation
quantiles_df = df.quantile([low, high])
print(quantiles_df)
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
0.0000	0.01	-3442.05957	-3492.263672	-3442.024658	
0.9973	60.00	0.00000	40.828285	0.000000	

	dropoff_latitude	passenger_count
0.0000	-3547.886719	0.0
0.9973	40.866947	6.0

**Step 2: Delete passengers more than 6.** We saw in the quantile plot above that 3 sigma as well as 3.5 sigma for passenger\_count is 6.0, so let's drop rows having passenger\_count > 6

```
In [10]: #Dropping rows
df = df[df.passenger_count <= 6]
df.shape
```

```
Out[10]: (55392419, 7)
```

Also, removing rows with 0 passengers as a taxi cannot operate with 0 passengers

```
In [0]: #Dropping rows
df = df[df.passenger_count > 0]
```

```
In [12]: df.shape
```

```
Out[12]: (55197444, 7)
```

**Step 3: Cleaning based on Coordinates** Now we try to understand latitudes and longitudes and remove outliers in these fields.

From basic Geography, we know that **Latitude** lies between -90 to +90. And **Longitude** lies between -180 to +180.

But in our observation, we have a longitude of **-3446**, which clearly means that we need to remove outliers like these.

Leaving everything aside, I found out the bounding box coordinates for NYC as the cabs were in NYC. Reducing our observations to these bounding box coordinates would ensure that we have eliminated all outliers.

The coordinates for New York City are: -74.259087, 40.477398, -73.70018, 40.91618

```
In [15]: df = df[(df.pickup_longitude >= -74.259087) & (df.dropoff_longitude <= -73.70018)
                & (df.dropoff_longitude >= -74.259087) & (df.pickup_longitude <= -73.70018)
                & (df.pickup_latitude >= 40.477398) & (df.dropoff_latitude <= 40.91618)
                & (df.dropoff_latitude >= 40.477398) & (df.pickup_latitude <= 40.91618)]
df.shape
```

```
Out[15]: (53979707, 7)
```

```
In [0]: df.describe()
```

Now we have sufficiently cleaned the data and it is free from outliers and looks uniform!

## 0.5 Creating a new feature for Time of Day

```
In [0]: df['time_of_day'] = df['pickup_datetime'].apply(lambda x:x.hour)
        #df.head()
```

## 0.6 Creating a new feature for Month of Year

```
In [0]: df['month_of_year'] = df['pickup_datetime'].apply(lambda x:x.month)
        #df.head()
```

## 0.7 Creating a new feature for Year

```
In [0]: df['year'] = df['pickup_datetime'].apply(lambda x:x.year)
        #df.head()
```

## 0.8 Creating a new feature for Day of Week

```
In [0]: df['day_of_week'] = pd.to_datetime(df['pickup_datetime']).dt.weekday.astype(np.int64)
```

```
In [0]: df.head()
```

```
In [26]: df.describe()
```

```
Out [26]:
```

	fare_amount	pickup_longitude	pickup_latitude	dropoff_longitude	\
count	5.397971e+07	5.397971e+07	5.397971e+07	5.397971e+07	
mean	1.126501e+01	-7.397079e+01	4.075009e+01	-7.397059e+01	
std	9.294146e+00	3.441382e-02	2.688210e-02	3.377897e-02	
min	2.000000e+00	-7.425903e+01	4.047759e+01	-7.425906e+01	
25%	6.000000e+00	-7.399229e+01	4.073658e+01	-7.399158e+01	
50%	8.500000e+00	-7.398210e+01	4.075336e+01	-7.398061e+01	
75%	1.250000e+01	-7.396835e+01	4.076752e+01	-7.396541e+01	
max	9.500000e+01	-7.370018e+01	4.091616e+01	-7.370018e+01	

	dropoff_latitude	passenger_count	time_of_day	month_of_year	\
count	5.397971e+07	5.397971e+07	5.397971e+07	5.397971e+07	
mean	4.075017e+01	1.691365e+00	1.351012e+01	6.269418e+00	
std	3.071281e-02	1.307223e+00	6.516493e+00	3.436574e+00	
min	4.047765e+01	1.000000e+00	0.000000e+00	1.000000e+00	
25%	4.073559e+01	1.000000e+00	9.000000e+00	3.000000e+00	
50%	4.075385e+01	1.000000e+00	1.400000e+01	6.000000e+00	
75%	4.076837e+01	2.000000e+00	1.900000e+01	9.000000e+00	
max	4.091618e+01	6.000000e+00	2.300000e+01	1.200000e+01	

	year	day_of_week
count	5.397971e+07	5.397971e+07
mean	2.011739e+03	3.041228e+00
std	1.865315e+00	1.949146e+00
min	2.009000e+03	0.000000e+00
25%	2.010000e+03	1.000000e+00
50%	2.012000e+03	3.000000e+00
75%	2.013000e+03	5.000000e+00
max	2.015000e+03	6.000000e+00

## 0.9 Calculating Euclidean Distance [Haversine Distance]

Referred this site: <https://stackoverflow.com/questions/27928/calculate-distance-between-two-latitude-longitude-points-haversine-formula?rq=1>

```
In [0]: from math import sin, cos, asin, sqrt, radians
```

```
def euclidean_dist(lon1, lat1, lon2, lat2):
    lon1, lat1, lon2, lat2 = map(np.radians, [lon1, lat1, lon2, lat2])
    # haversine formula
    dlon = lon2 - lon1
    dlat = lat2 - lat1
    a = (np.sin(dlat/2)**2
          + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2)**2)
    c = 2 * np.arcsin(np.sqrt(a))
    miles = 3959 * c
    return miles
```

```
In [0]: df['euclidean_dist'] = euclidean_dist(df.loc[:, 'pickup_longitude'].astype(float),
                                              df.loc[:, 'pickup_latitude'].astype(float),
                                              df.loc[:, 'dropoff_longitude'].astype(float),
                                              df.loc[:, 'dropoff_latitude'].astype(float))
```

```
In [0]: df.head()
```

```
Out[0]:
```

	fare_amount		pickup_datetime	pickup_longitude	pickup_latitude \
0	4.5	2009-06-15	17:26:00+00:00	-73.844315	40.721317
1	16.9	2010-01-05	16:52:00+00:00	-74.016045	40.711304
2	5.7	2011-08-18	00:35:00+00:00	-73.982735	40.761269
3	7.7	2012-04-21	04:30:00+00:00	-73.987129	40.733143
4	5.3	2010-03-09	07:51:00+00:00	-73.968094	40.768009

	dropoff_longitude	dropoff_latitude	passenger_count	time_of_day \
0	-73.841614	40.712276	1	17
1	-73.979271	40.782005	1	16
2	-73.991241	40.750561	2	0
3	-73.991570	40.758091	1	4
4	-73.956657	40.783764	1	7

	month_of_year	year	euclidean_dist
0	6	2009	0.640513
1	1	2010	5.250911
2	8	2011	0.863531
3	4	2012	1.739456
4	3	2010	1.242248

```
In [0]: df.describe()
```

### 0.9.1 Data Cleaning after adding new features

After calculating Euclidean distance, there is an important parameter to work with. We now will **clean** the rows that have distance travelled less than 0.3 miles as most of the people do not use cabs for distances as short as this! There are almost 2 million such entries and they can spoil the model

```
In [22]: len(df[df.euclidean_dist <= 0.3 ])
```



```
Out[22]: 2093705
```

```
In [0]: #Removing rows having distance less than 0.3 miles
df = df[df.euclidean_dist >= 0.3]
```

```
In [24]: df.shape
```

```
Out[24]: (51886002, 12)
```

## 1 Now let's export this cleaned file!

```
In [0]: #Creating a feather file
import feather
feather.write_dataframe(df, 'gdrive/My Drive/Colab Notebooks/Datasets/NYC/train_cleaned')
```

### 1.1 Creating a correlation DataFrame

```
In [0]: df_correlation = df.filter(['fare_amount', 'time_of_day', 'euclidean_dist'], axis=1)
```

## 2 Calculating Pearson Correlation Coefficient

```
In [27]: df_correlation.corr(method='pearson')
```

```
Out[27]:
```

	fare_amount	time_of_day	euclidean_dist
fare_amount	1.000000	-0.018592	0.904927
time_of_day	-0.018592	1.000000	-0.032418
euclidean_dist	0.904927	-0.032418	1.000000

```
In [0]: sns.pairplot(df_correlation)
```

### 2.1 Converting to CSV File

```
In [0]: #df.to_csv(path="gdrive/My Drive/Colab Notebooks/Datasets/NYC/train_cleaned.csv")
df.to_csv("gdrive/My Drive/Colab Notebooks/Datasets/NYC/train_cleaned.csv", encoding='utf-8')
```

### 2.2 Converting to Feather File

```
In [0]: df_1 = df
df_1.reset_index()
#df_1.head()
#df_1.to_feather("gdrive/My Drive/Colab Notebooks/Datasets/NYC/train_cleaned.feather")
df_1.to_feather("gdrive/My Drive/Colab Notebooks/Datasets/NYC/train_cleaned.feather")
```

```
Out[0]:
```

	fare_amount		pickup_datetime	pickup_longitude	pickup_latitude	\
0	4.5	2009-06-15	17:26:00+00:00	-73.844315	40.721317	
1	16.9	2010-01-05	16:52:00+00:00	-74.016045	40.711304	
2	5.7	2011-08-18	00:35:00+00:00	-73.982735	40.761269	
3	7.7	2012-04-21	04:30:00+00:00	-73.987129	40.733143	

4	5.3	2010-03-09 07:51:00+00:00	-73.968094	40.768009
---	-----	---------------------------	------------	-----------

	dropoff_longitude	dropoff_latitude	passenger_count	time_of_day \
0	-73.841614	40.712276	1	17
1	-73.979271	40.782005	1	16
2	-73.991241	40.750561	2	0
3	-73.991570	40.758091	1	4
4	-73.956657	40.783764	1	7


  

	month_of_year	year	euclidean_dist
0	6	2009	0.640513
1	1	2010	5.250911
2	8	2011	0.863531
3	4	2012	1.739456
4	3	2010	1.242248

## 2.3 Plotting the relationship between variables

```
In [0]: sns.pairplot(df_correlation[:1000000])
```

```
Out[0]: <seaborn.axisgrid.PairGrid at 0x7f23831a6198>
```




The plots do not show a very linear kind of relationship, but fare\_amount and euclidean\_dist are highly correlated

```
In [0]: df_reduced = df.sample(frac=0.2, random_state=42)
```

```
In [31]: #Creating a distribution plot to check the frequency of cab bookings according to the
sns.distplot(df_reduced['time_of_day'], color="green")
```


```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x7fac4b58a780>
```



From this graph, it is evident that the highest number of cab bookings are done in between 6PM to 9PM, and this coincides with the closing time of offices in NYC. Also the bookings done at late night [12AM to 5AM] could be people travelling to/from airports!

```
In [30]: sns.kdeplot(df_reduced['day_of_week'])
```

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7fac6c09a710>
```




This plot shows that the distribution is more or less similar for every day I have drawn a Kernel Density Estimation plot here!

```
In [0]: sns.jointplot("euclidean_dist", "fare_amount", data=df_reduced, kind='hex')
```

```
In [0]: df_reduced = df_reduced[:1000000]
```

```
In [35]: sns.jointplot("euclidean_dist", "fare_amount", data=df_reduced,color='green', kind='r
```

```
Out[35]: <seaborn.axisgrid.JointGrid at 0x7fac4b36b7b8>
```



This plot also reveals an interesting insight into the data! The correlation is quite high and can be seen from the line!

### 3 Building Linear Regression Model based on the above features

```
In [0]: import numpy as np           #For NumPy arrays and mathematical functions
import pandas as pd                 #For operations using DataFrames
import matplotlib.pyplot as plt     #For plotting
import seaborn as sns               #For plotting
import time
import os
```

```
In [2]: dir_path = "gdrive/My Drive/Colab Notebooks/Datasets/NYC/"
file_name = "train_cleaned.feather"
```

```

df_lin = pd.read_feather(dir_path+file_name, nthreads=2)
%time

/usr/local/lib/python3.6/dist-packages/pandas/io/feather_format.py:112: FutureWarning: `nthreads`
return feather.read_dataframe(path, nthreads=nthreads)

CPU times: user 1e+03 ns, sys: 2 µs, total: 3 µs
Wall time: 9.78 µs

In [0]: print(df_lin.shape)
df_lin.head()

In [0]: from sklearn.linear_model import LinearRegression
from sklearn.cross_validation import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import PolynomialFeatures
#from IPython.html import widgets
#from IPython.html.widgets import interact
#from IPython.display import display

In [0]: df_subset = df_lin.sample(frac=0.5, random_state=42)

In [14]: df_subset.shape

Out[14]: (25943001, 12)

In [0]: features = ['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude', 'fare_amount']

X1 = df_subset[features].values

Y = df_subset[['fare_amount']].values

X1_train, X1_test, y_train, y_test = train_test_split(X1, Y, test_size=0.3, random_state=42)

In [16]: slr = LinearRegression()
slr.fit(X1_train, y_train)
y_train_pred = slr.predict(X1_train)
y_test_pred = slr.predict(X1_test)
slr.coef_

Out[16]: array([[ 1.26752508e+01,  8.25524376e+00, -5.74357322e+00,
-1.50032256e+01,  3.70326356e+00,  1.07440275e-02,
-4.96219671e-02]])

When frac=0.5 array([[ 1.26752508e+01,  8.25524376e+00, -5.74357322e+00, -1.50032256e+01,
3.70326356e+00, 1.07440275e-02, -4.96219671e-02]])
For frac=0.2array([[ 1.26912264e+01,  8.28425098e+00, -5.60640618e+00, -1.50239738e+01,
3.70220438e+00, 1.05555875e-02, -5.09145992e-02]])

```

```
In [17]: from math import sqrt
import numpy as np
print ('Simple Linear Regression on Euclidean Distance Results:')
print ('Root Mean Squared Error Train: %.5f' % np.sqrt(mean_squared_error(y_train,y_test)))
print ('Root Mean Squared Error Test: %.5f' % np.sqrt(mean_squared_error(y_test,y_test)))
```

Simple Linear Regression on Euclidean Distance Results:

Root Mean Squared Error Train: 3.87634

Root Mean Squared Error Test: 3.87814

Simple Linear Regression over 50% of dataset Results: Root Mean Squared Error Train: 3.87634  
Root Mean Squared Error Test: 3.87814

Simple Linear Regression frac=0.2 Results: Root Mean Squared Error Train: 3.86948 Root Mean Squared Error Test: 3.87446

## 4 Checking results with the Test File

```
In [0]: dir_path = "gdrive/My Drive/Colab Notebooks/Datasets/NYC/"
file_name = "test.csv"
```

```
In [0]: df_test = pd.read_csv(dir_path+file_name)
```

```
In [0]: df_test.head()
```

```
In [0]: from math import sin, cos, asin, sqrt, radians
```

```
def euclidean_dist(lon1, lat1, lon2, lat2):
    lon1, lat1, lon2, lat2 = map(np.radians, [lon1, lat1, lon2, lat2])
    # haversine formula
    dlon = lon2 - lon1
    dlat = lat2 - lat1
    a = (np.sin(dlat/2)**2
          + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2)**2)
    c = 2 * np.arcsin(np.sqrt(a))
    miles = 3959 * c
    return miles
```

```
In [0]: df_test['euclidean_dist'] = euclidean_dist(df_test.loc[:, 'pickup_longitude'].astype(float),
                                                    df_test.loc[:, 'pickup_latitude'].astype(float),
                                                    df_test.loc[:, 'dropoff_longitude'].astype(float),
                                                    df_test.loc[:, 'dropoff_latitude'].astype(float))
```

```
In [0]: df_test['time_of_day'] = pd.to_datetime(df_test['pickup_datetime']).dt.hour
df_test['year'] = pd.to_datetime(df_test['pickup_datetime']).dt.year
df_test['day_of_week'] = pd.to_datetime(df_test['pickup_datetime']).dt.weekday.astype(int)
df_test['month_of_year'] = pd.to_datetime(df_test['pickup_datetime']).dt.month
```

```
In [10]: df_test.shape
```



```
Out[10]: (9914, 12)
```

```
In [11]: features = ['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude']
```

```
X1_train = df_subset[features].values
y_train = df_subset[['fare_amount']].values
X1_test = df_test[features].values
```

```
slr = LinearRegression()
model = slr.fit(X1_train, y_train)
pred = model.predict(X1_test)
slr.coef_
```

```
print(pred[:10])
```

```
[[ 9.23534548]
 [ 9.10376691]
 [ 4.82819657]
 [ 8.15740108]
 [16.50943335]
 [11.40557623]
 [ 5.77304848]
 [54.6477168 ]
 [11.84817396]
 [ 6.09650996]]
```

```
In [12]: df_submit = df_test[['key']]
df_submit['fare_amount'] = pred
df_submit.to_csv(dir_path+"linear_reg_2.csv", index=False)
print("Done!")
```

Done!

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

```
In [0]: df_submit.shape
```

```
Out[0]: (9914, 2)
```

#### 4.1 Kaggle RMSE for Linear Regression: 5.814, 5.912, 5.813

## 5 Random Forest

```
In [0]: import numpy as np                #For NumPy arrays and mathematical functions
import pandas as pd                      #For operations using DataFrames
import matplotlib.pyplot as plt          #For plotting
import seaborn as sns                    #For plotting
import time
import os

In [0]: from google.colab import drive
drive.mount('gdrive')

In [0]: dir_path = "gdrive/My Drive/Colab Notebooks/Datasets/NYC/"
file_name = "train_cleaned.feather"

In [0]: !pip install -U feather-format

In [0]: df = pd.read_feather(dir_path+file_name, nthreads=2)

/usr/local/lib/python3.6/dist-packages/pandas/io/feather_format.py:112: FutureWarning: `nthreads`
return feather.read_dataframe(path, nthreads=nthreads)

In [0]: # Loading 20% of the dataset
df_subset = df.sample(frac=0.1, random_state=42)
df_subset.head()
```

Out[0]:

	fare_amount	pickup_datetime	pickup_longitude	\
44976798	12.5	2011-02-15 11:29:00+00:00	-73.998558	
45538027	4.9	2009-09-26 04:04:00+00:00	-73.988586	
47711137	10.1	2012-07-28 10:09:00+00:00	-73.992920	
19321254	6.5	2012-05-09 22:24:00+00:00	-73.986252	
49976425	12.5	2012-02-08 20:05:00+00:00	-73.994736	

	pickup_latitude	dropoff_longitude	dropoff_latitude	\
44976798	40.724583	-73.995834	40.767811	
45538027	40.733788	-73.976578	40.727203	
47711137	40.755295	-73.977371	40.784374	
19321254	40.726101	-73.981277	40.744148	
49976425	40.726070	-73.966980	40.758049	

	passenger_count	time_of_day	month_of_year	year	euclidean_dist	\
44976798	1	11	2	2011	2.990362	
45538027	1	4	9	2009	0.776115	
47711137	4	10	7	2012	2.167814	
19321254	6	22	5	2012	1.273934	
49976425	1	20	2	2012	2.644608	

day\_of\_week

44976798	1
45538027	5
47711137	5
19321254	2
49976425	2

```
In [0]: df_subset.shape
```

```
Out[0]: (518860, 12)
```

```
In [0]: features = ['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude',
                    'passenger_count', 'euclidean_dist', 'year', 'day_of_week', 'time_of_day']
X = df_subset[features].values
y = df_subset['fare_amount'].values
```

```
In [0]: # Using Skicit-learn to split data into training and testing sets
from sklearn.model_selection import train_test_split
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, random_state = 42)
```

```
In [0]: # Import the model we are using
from sklearn.ensemble import RandomForestRegressor
# Instantiate model with 10 decision trees
rf = RandomForestRegressor(n_estimators = 50, random_state = 42)
# Train the model on training data
rf.fit(X_train, y_train)
y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)
```

```
In [0]: from math import sqrt
import numpy as np
from sklearn.metrics import mean_squared_error
print ('Random Forest Results:')
print ('Root Mean Squared Error Train: %.5f' % np.sqrt(mean_squared_error(y_train, y_train_pred)))
print ('Root Mean Squared Error Test: %.5f' % np.sqrt(mean_squared_error(y_test, y_test_pred)))
```

## 6 Checking with Test File

```
In [0]: dir_path = "gdrive/My Drive/Colab Notebooks/Datasets/NYC/"
file_name = "test.csv"
```

```
In [0]: df_test = pd.read_csv(dir_path+file_name)
```

```
In [0]: from math import sin, cos, asin, sqrt, radians
```

```
def euclidean_dist(lon1, lat1, lon2, lat2):
    lon1, lat1, lon2, lat2 = map(np.radians, [lon1, lat1, lon2, lat2])
    # haversine formula
    dlon = lon2 - lon1
```

```

    dlat = lat2 - lat1
    a = (np.sin(dlat/2)**2
          + np.cos(lat1) * np.cos(lat2) * np.sin(dlon/2)**2)
    c = 2 * np.arcsin(np.sqrt(a))
    miles = 3959 * c
    return miles

In [0]: df_test['euclidean_dist'] = euclidean_dist(df_test.loc[:, 'pickup_longitude'].astype(float),
                                                    df_test.loc[:, 'pickup_latitude'].astype(float),
                                                    df_test.loc[:, 'dropoff_longitude'].astype(float),
                                                    df_test.loc[:, 'dropoff_latitude'].astype(float))

In [0]: df_test['time_of_day'] = pd.to_datetime(df_test['pickup_datetime']).dt.hour
df_test['year'] = pd.to_datetime(df_test['pickup_datetime']).dt.year
df_test['day_of_week'] = pd.to_datetime(df_test['pickup_datetime']).dt.weekday.astype(int)
df_test['month_of_year'] = pd.to_datetime(df_test['pickup_datetime']).dt.month

In [0]: df_test.head()

In [0]: features = ['pickup_longitude', 'pickup_latitude', 'dropoff_longitude', 'dropoff_latitude',
                    'passenger_count', 'euclidean_dist', 'year', 'day_of_week', 'time_of_day']
X_train = df_subset[features].values
y_train = df_subset['fare_amount'].values
X_test = df_test[features].values

In [0]: # Import the model we are using
from sklearn.ensemble import RandomForestRegressor
# Instantiate model with 10 decision trees
rf = RandomForestRegressor(n_estimators = 50, random_state = 42)
# Train the model on training data
rf.fit(X_train, y_train)
y_train_pred = rf.predict(X_train)
y_test_pred = rf.predict(X_test)

print(y_test_pred[:10])

In [0]: from math import sqrt
import numpy as np
from sklearn.metrics import mean_squared_error
print ('Random Forest Results:')
print ('Root Mean Squared Error Train: %.5f' % np.sqrt(mean_squared_error(y_train, y_train_pred)))
print ('Root Mean Squared Error Test: %.5f' % np.sqrt(mean_squared_error(y_test, y_test_pred)))

In [0]: df_submit = df_test[['key']]
df_submit['fare_amount'] = y_test_pred
df_submit.to_csv(dir_path+"rf_regressor_n100.csv", index=False)
print("Done!")

```

Done!

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html>

## 7 Random Forest RMSE over Kaggle:

7.1 1. No. of rows =  $1/100$  \* Cleaned Dataset [Around 520,000 rows] and `n_estimators` = 50 -> *RMSE*: 3.45618

7.2 2. No. of rows =  $1/10$  \* Cleaned Dataset [Around 5.2 million rows] and `n_estimators` = 100 -> *RMSE*: 3.43151

7.3 END OF FILE

---